

# 17-614: Project Write-up: Modeling *Presto* (Ride Sharing App)

Team #12 | Iris Huang, Viren Dodia, Ziqin Shen, Ray Xue

October 5, 2025

# 1 Task 1: Structural Modeling

## 1.1 Object Model Diagram

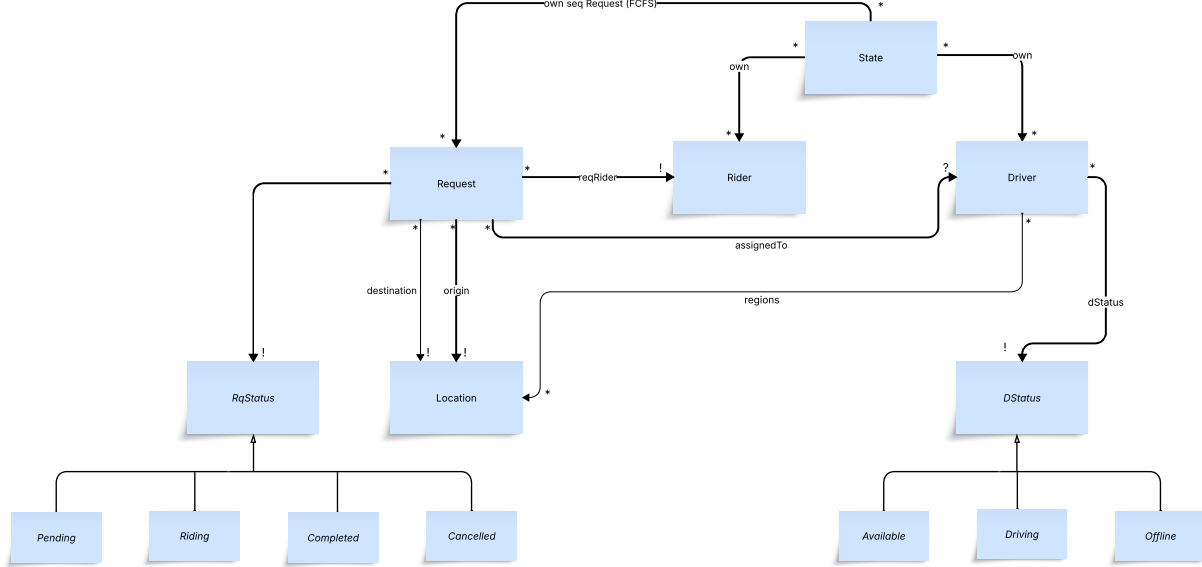


Figure 1: Object model of the *Presto* system.

## 1.2 Invariants Discovered During Modeling

During Alloy modeling of *Presto*, we identified several invariants that ensure consistency of the system state:

- **One active request per rider:** Each rider can have at most one request that is either **Pending** or **Riding**.
- **Driver exclusivity:** Each driver can serve at most one request at a time. A driver is in the **Driving** state if and only if they are assigned to exactly one active request.
- **Assignment consistency:** A request in **Riding** must have exactly one assigned driver. Requests in **Pending**, **Completed**, or **Cancelled** must have no assigned driver.
- **Queue well-formedness:** The set of requests in the **Pending** state must exactly equal the elements of the `pendingQ` sequence (with no duplicates).
- **Origin-destination sanity:** For realism, each request must have distinct origin and destination.

These invariants capture both explicit rules in the specification and implicit requirements necessary to preserve logical correctness.

### 1.3 Model Validation Strategy

Our validation strategy for the Alloy model involved:

- **Operation preservation checks:** We wrote assertions ensuring that the four core operations (`request`, `cancel`, `match`, and `complete`) preserve the system invariants across state transitions.
- **Visualization of states:** Using `run` commands, we generated concrete states, such as an empty system, one pending request, and one riding request. These helped confirm intuitive behavior.
- **Counterexample analysis:** If Alloy produced a counterexample, we refined the model (e.g., corrected scoping errors or missing conditions) until the invariants held.
- **Cross-checking with the spec:** Each invariant and operation was traced back to requirements in the Presto specification to ensure coverage.

In addition, we developed a suite of **positive and negative test predicates** to exercise the model further. Positive tests (e.g., `test_MultiplePending`, `test_MultipleConcurrentRides`) confirmed that realistic states were satisfiable, while negative tests (e.g., `test_RidingRequestInPendingQueue`, `test_AvailableDriverIsAssigned`, `test_DuplicateRequestInQueue`) demonstrated that invalid states were impossible under our invariants. Running these tests under scopes of 4–7 atoms with exactly one `State` reinforced our confidence that the model was both correct and complete.

### 1.4 Scopes for Checking Assertions

For invariant preservation checks, we used scopes of 6–7 objects with exactly one `State` and up to 6 sequence elements. This bound was sufficient because:

- It allowed us to explore diverse combinations of Riders, Drivers, and Requests while keeping the analysis tractable.
- Larger scopes (beyond 7) did not uncover new counterexamples, suggesting our invariants are robust.
- These bounds match the project’s guidance for balancing coverage with solver performance.

## 2 Task 2: Concurrency with FSP/LTSA

### 2.1 Process Structure Diagram

- **ORDER:** a generator that cycles order identifiers  $t \in \{1, \dots, \text{NUMRIDER}\}$  and emits `request[t]` in round-robin order.
- **RIDER(*i*)** for each  $i \in \{1, \dots, \text{NUMRIDER}\}$ : `request[i] → match[i] → complete[i]` (or `cancel[i]` before `match`).
- **SCHEDULER** as a driver pool `DRIVER(d, t)`:
  - if  $d > 0$ : `match[t]` uses one available driver  $\Rightarrow \text{DRIVER}(d-1, \text{next}(t))$
  - `cancel[t]` skips the slot  $\Rightarrow \text{DRIVER}(d, \text{next}(t))$
  - if  $d < \text{NUMDRIVER}$ : `complete[t]` returns a driver  $\Rightarrow \text{DRIVER}(d+1, t)$

The system composition is:

$$\text{ASSIGN\_SYS} = (\parallel i : \text{Order} @ \text{RIDER}(i)) \parallel \text{SCHEDULER} \parallel \text{ORDER}.$$

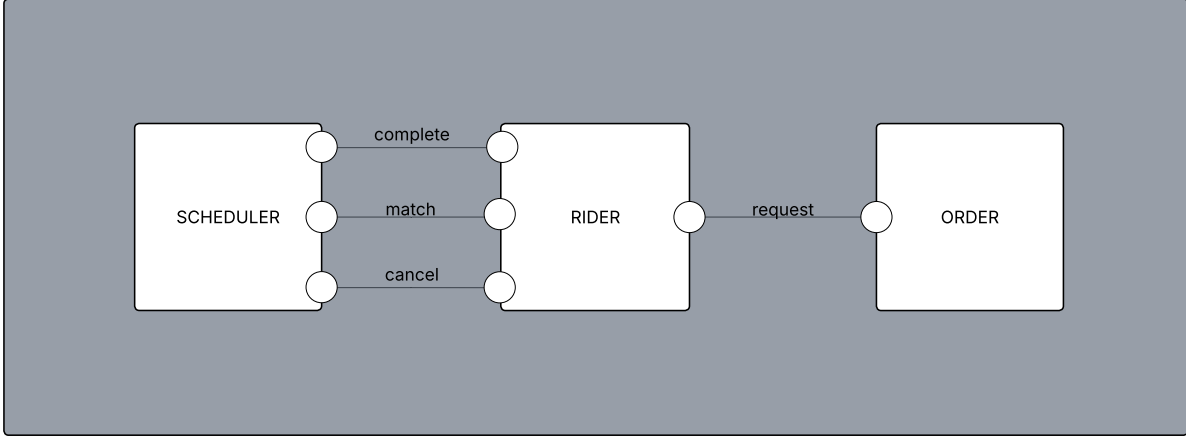


Figure 2: Process structure of the *Presto* system.

## 2.2 Protocol Design

**Events.** We model four observable events with an index  $t$  identifying the request/rider: `request`[ $t$ ], `match`[ $t$ ], `cancel`[ $t$ ], `complete`[ $t$ ].

**Rider protocol.** For each rider  $i$ , the process `RIDER`( $i$ ) performs `request`[ $i$ ] then waits either to be matched (`match`[ $i$ ]) and subsequently completed (`complete`[ $i$ ]), or to cancel before being matched (`cancel`[ $i$ ]), after which it may request again.

**Scheduler / driver pool.** The scheduler maintains an integer  $d$  of available drivers and a rotating pointer  $t$ . If  $d > 0$ , the scheduler may synchronize on `match`[ $t$ ] (consuming one driver and moving to `next`( $t$ )); upon `cancel`[ $t$ ] it skips to `next`( $t$ ) without consuming a driver; upon `complete`[ $t$ ], it returns a driver ( $d+1$ ) while keeping  $t$  to allow back-to-back completions.

**Arrival order (FCFS).** Arrival order is modeled by `ORDER` which emits `request`[ $t$ ] in a round-robin over the index set use a ticketing scheme to make sure the riders are served by requesting order, giving each user request a number to manage request order. We enforce first-come, first-served with a property automaton `FIFO_CHECK` that flags an `ERROR` if `match`[2] occurs before `match`[1] when 1 arrived first (and symmetrically for two requests). This is the standard FIFO witness for two requests and is sufficient to catch out-of-order matches.

## 2.3 Details Abstracted from Task 1

Compared to the Alloy state model, the FSP protocol abstracts away:

- **Locations and regions:** We do not carry `origin`, `destination`, or driver `regions`. Region feasibility is left implicit (always matchable when a driver is available).
- **Structural associations:** Fields such as `reqRider`, `assignedTo`, and the explicit `pendingQ` structure are represented behaviorally via indexed events and the scheduler's pointer  $t$ .
- **Per-request object identity:** Requests are denoted by indices ( $t$ ) rather than explicit objects; this is sufficient for the concurrency protocol.

## 2.4 Details Added Beyond Task 1

The FSP model adds explicit behavioral constraints that Alloy does not natively express:

- **Interleaving semantics:** Competing `request`, `match`, `cancel`, `complete` events from multiple riders interleave under LTSA’s process-algebra rules.
- **Driver capacity dynamics:** The integer parameter  $d$  captures resource consumption at `match[t]` and release at `complete[t]`.
- **Safety property for FCFS:** The `FIFO_CHECK` property process rules out any trace where a later request is matched before an earlier one (two-request witness).
- **Progress (liveness):** We require that completions keep happening via the LTSA progress set `progress SERVED = {complete[Order]}`; optionally, we can strengthen liveness with a per-request watchdog ensuring that after `request[i]` some `match[i]` or `cancel[i]` must occur.

## 3 Task 3: Reflection

### 3.1 Alloy: Strengths and Weaknesses

**Strengths:**

- Naturally suited for modeling structural constraints and invariants.
- Immediate visualization of counterexamples, which aids debugging.
- Compact and expressive syntax for relational properties.

**Weaknesses:**

- Not designed to capture concurrency or event ordering explicitly.
- Large scopes can cause performance issues.
- Expressing temporal behaviors (e.g., eventuality) requires workarounds.

### 3.2 FSP/LTSA: Strengths and Weaknesses

**Strengths.**

- Natural description of concurrent event interleavings and shared synchronizations.
- Built-in checks for deadlock and support for progress (liveness) via progress sets.
- Property processes (e.g., `FIFO_CHECK`) make safety conditions executable.

**Weaknesses.**

- Structural constraints (e.g., object associations, multiplicities) are not explicit; they must be encoded behaviorally or left to Alloy.
- FIFO beyond the two-request witness requires either a generalized property pattern or a larger monitor, which can grow in complexity.
- Liveness per request (“each request eventually served”) often needs additional fairness assumptions or watchdog processes to exclude pathological schedules.

### 3.3 Other Aspects of Ride Sharing

While our models capture the core protocol, real-world ride sharing involves additional aspects:

- **Pricing and payment:** Fare calculation, dynamic pricing, and payment handling.
- **Trust and reputation:** Ratings, cancellation penalties, and fraud prevention.
- **Geographic constraints:** Real-world routing, travel times, and multi-region rides.
- **System resilience:** Handling driver disconnections, rider no-shows, or sudden surges.

Additional information : We used AI(ChatGPT 5) on Alloy, FSP modeling and latex formating